

The NIX file format for data and metadata: Acquire together, file together


Adrian Stoewer
Christian Kellner

Introduction

- Why another file format?
 - Many formats are proprietary and not open source
 - Often poorly documented
 - Existing open source formats do not support complex metadata
- About the NIX project
 - Started at the hackathon after the 2012 INCF Congress
 - Project goals:
 - A flexible format for neuroscience data and metadata based on HDF5
 - Development of a data model that can also be used in other back-ends
 - Development of a reference implementation in C++

Talking about data:
what do we need to store?


Sampled Data



What do we need to store?

- Name
- 1D Data
 $[s_1, \dots, s_n]$
- label (y-axis)
- unit (y-axis)
- label (x-axis)
- sampling unit (x-axis)
- sampling interval (x-axis)

Irregularly Sampled Data



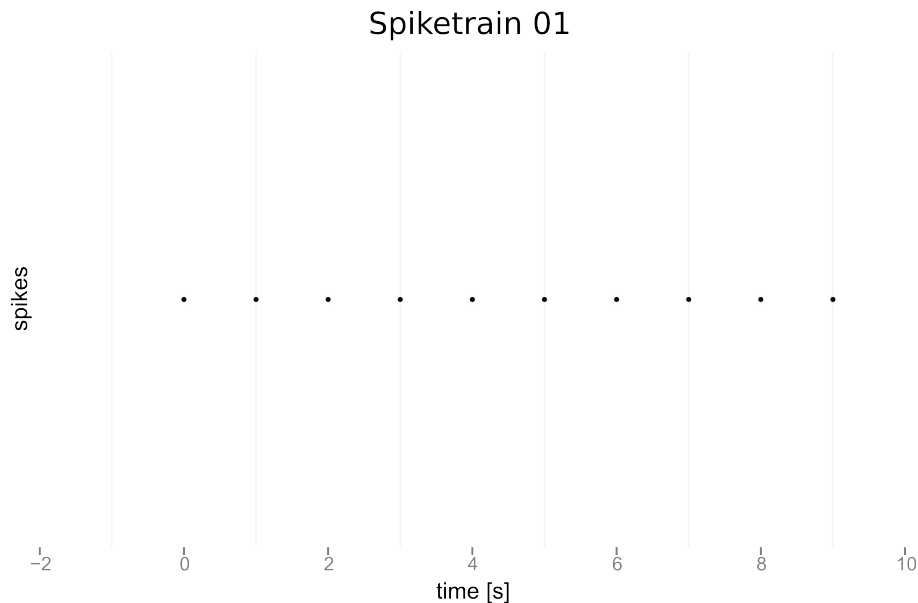
What do we need to store?

- Name
- 1D Data
 $[s_1, \dots, s_n]$
- 1D Data (ticks)
 $[t_1, \dots, t_n]$
- label (y-axis)
- unit (y-axis)
- label (x-axis)
- unit (x-axis)

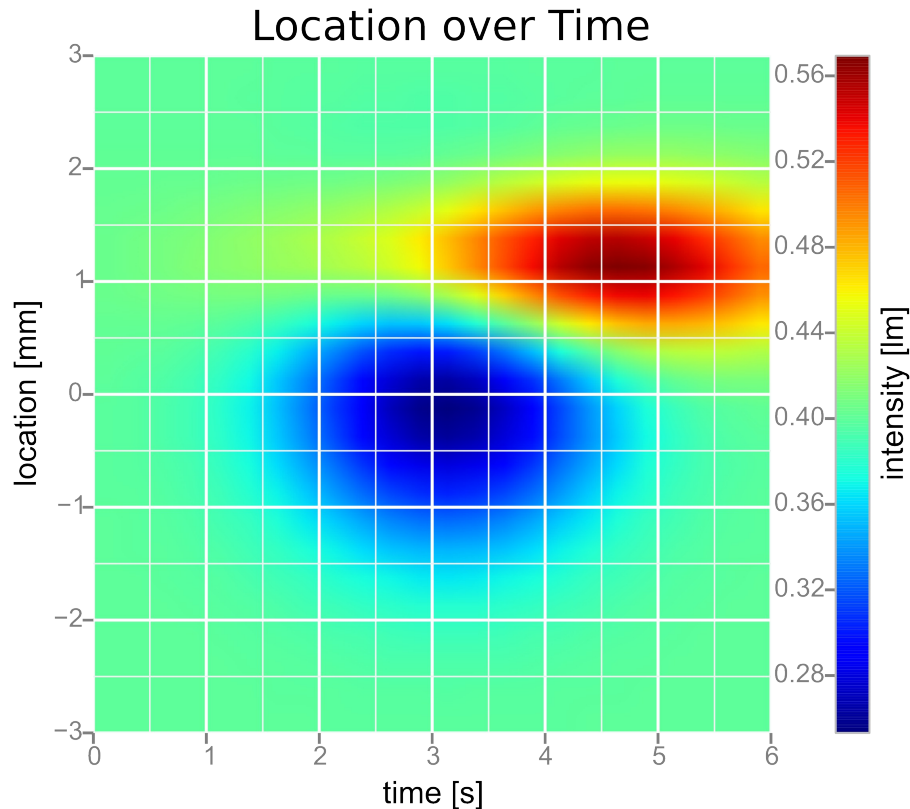
Event Data

What do we need to store?

- Name
- 1D Data
 $[e_1, \dots, e_n]$
- label (y-axis)
- label (x-axis)
- unit (x-axis)



Multiple Dimensions



What do we need to store?


- Name
- 2D Data
 $[[s_{1,1}, \dots, s_{n,m}]]$
- label (z-axis)
- unit (z-axis)
- label (x- / y-axis)
- sampling unit (x- / y-axis)
- sampling interval (x- / y-axis)

Data Models

Data Models

Why is a data model is useful?

- Allows early evaluation of different use cases
- Captures the most important concepts of the format
- Format specification can be derived from the model
- API or programming model can be derived from the model



Conventional Data Models

- Data model development process
 - Isolation of data objects the model should be able to describe
 - Define entities reflecting the data objects
 - Define properties and relationships of the entities
- Advantages
 - Model is easy to understand
 - Entities have a strong semantic meaning
- Disadvantages
 - Can result in a very large number of entities

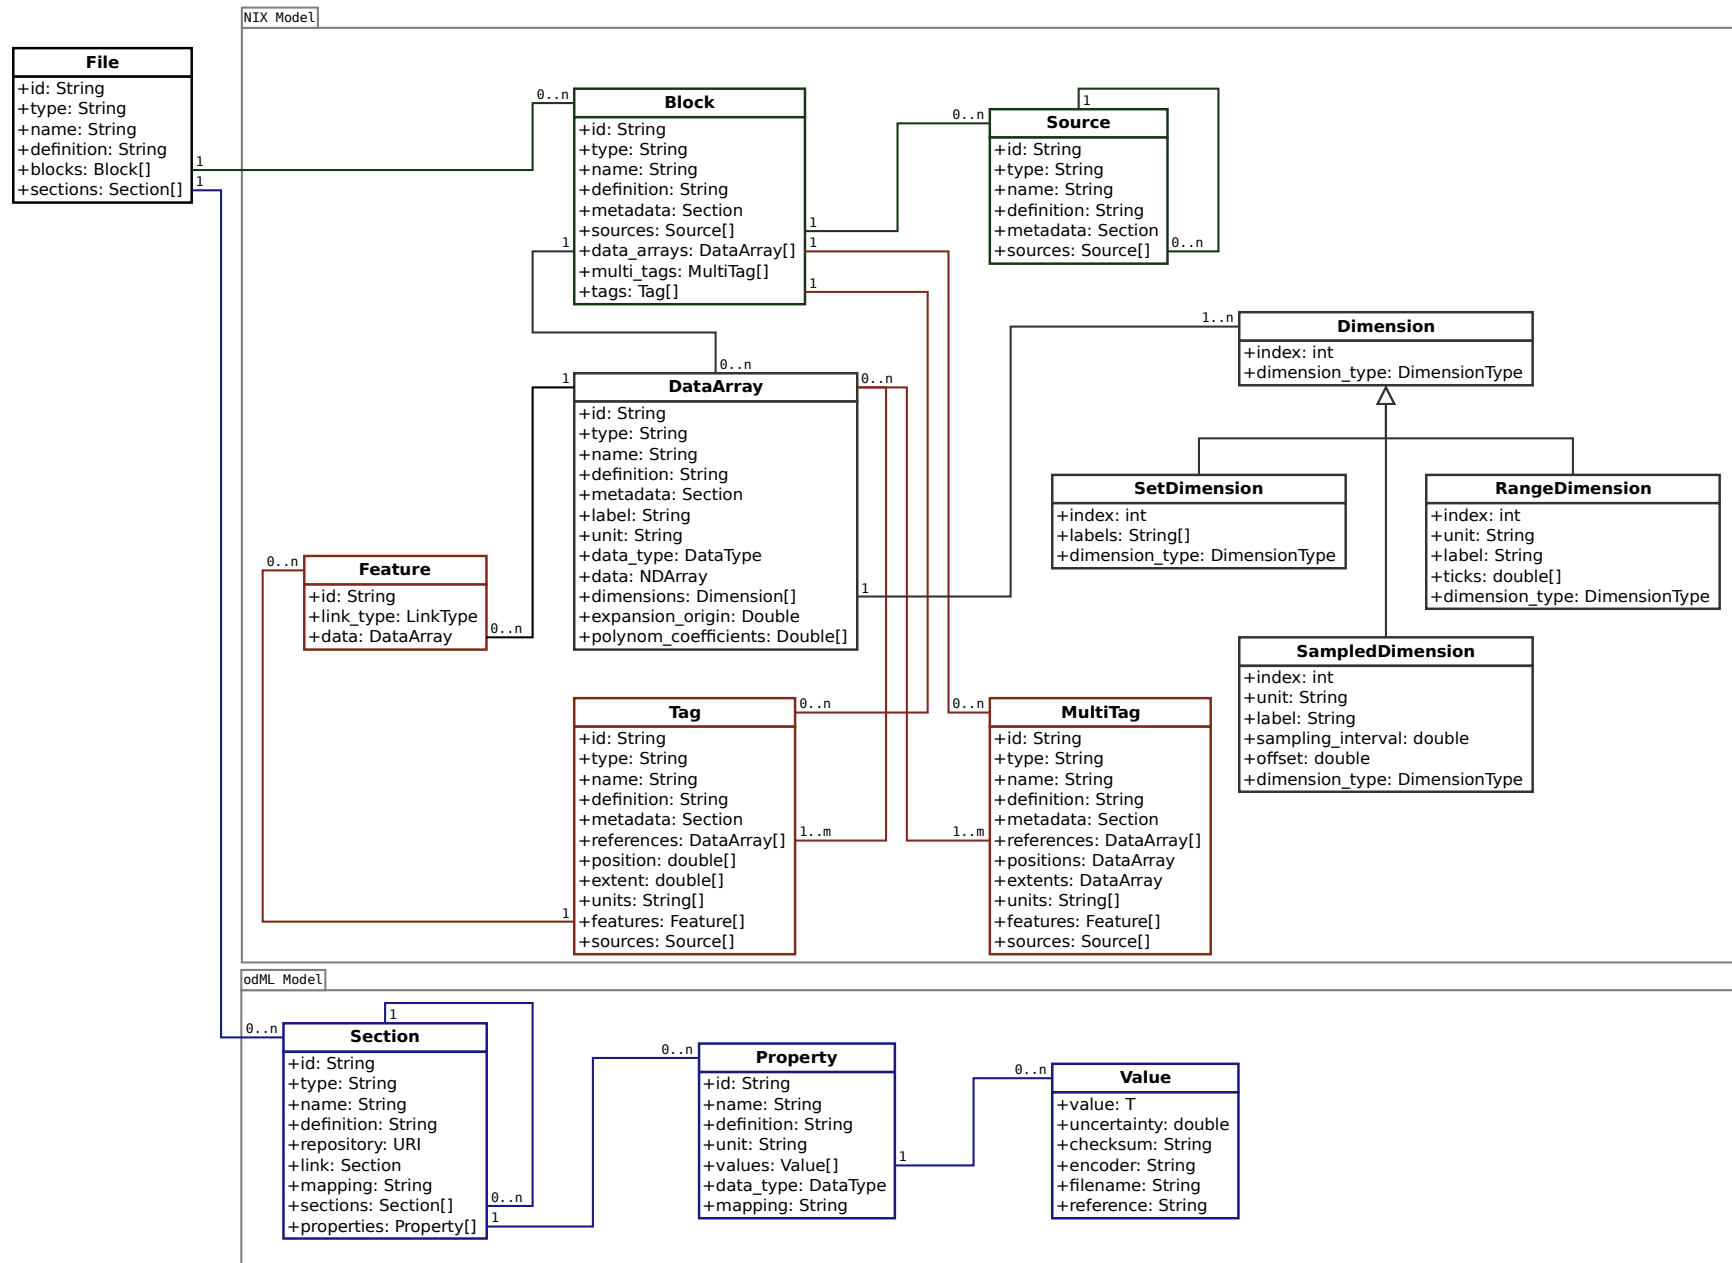
Generic Data Models

- Data model development process
 - Isolation of data objects the model should be able to describe
 - Identify common properties and concepts of data objects
 - Design generic entities that can represent those common concepts
 - Introduce generic properties that can describe differences between data objects
- Advantages
 - Model is much smaller
- Disadvantages
 - May be harder to understand and to process
 - Entities have less semantic meaning

The NIX approach

- Model for data
 - Take entities from existing formats such as NEO
 - Extract common concepts from these entities
 - Generalize the concepts as far as possible
- Model for metadata
 - Use the odML data model
 - Connect both data models

NIX Model: Overview



NIX Model: Common Entity Properties

- id:
 - Randomly generated uuid
 - Low collision probability
- name:
 - Human readable user defined identifier
- type:
 - Provides domain specificity to the generic entities
 - The type may be defined in a terminology or ontology
- definition:
 - User defined description of the entity

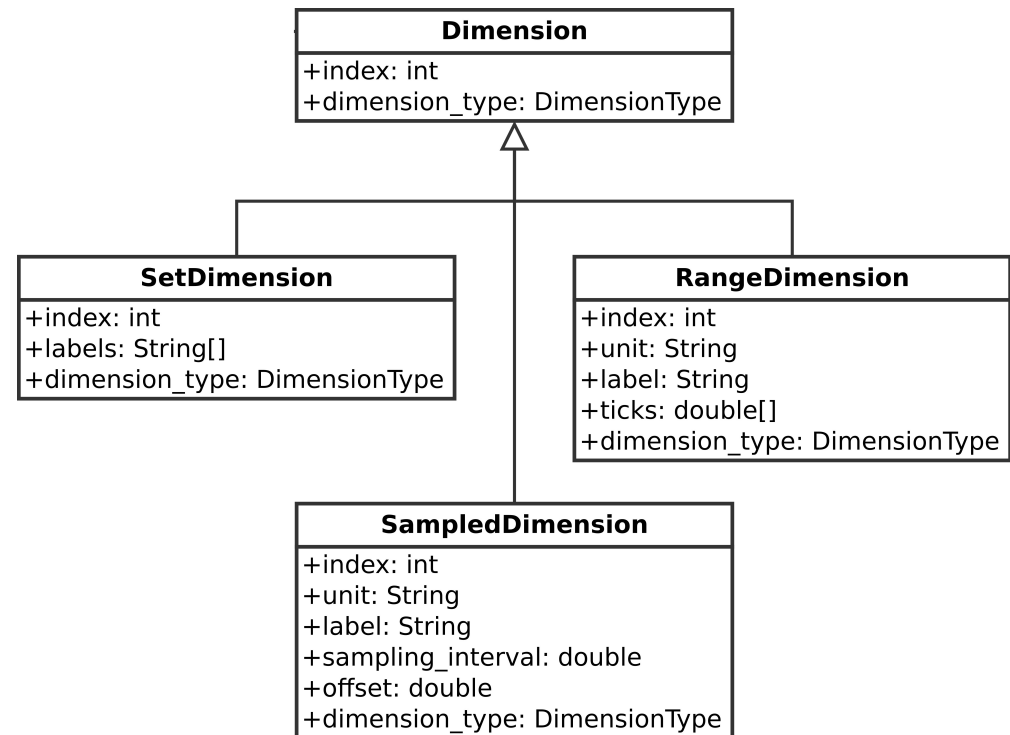
NIX Model: DataArray

- Main entity for storing data
- Stores data in n-dimensional array
- Provides unit, data type and label for the stored values
- Has a dimension descriptor for each dimension of the data array
- Additional metadata via link to an odML section
- Value scaling

DataArray
+id: String +type: String +name: String +definition: String +metadata: Section +label: String +unit: String +data_type: DataType +data: NDAarray +dimensions: Dimension[] +expansion_origin: Double +polynom_coefficients: Double[]

NIX Model: Dimensions

- Describe the dimensions of data in a DataArray entity
- Three different kinds of dimensions
 - Sampled
 - Range
 - Set



NIX Model: Tags

- Can define points or regions of interest
- Examples
 - Events
 - Spikes
 - Epochs
- MultiTag can use DataArray entities to define positions and extents

Tag
+id: String +type: String +name: String +definition: String +metadata: Section +references: DataArray[] +position: double[] +extent: double[] +units: String[] +features: Feature[] +sources: Source[]

MultiTag
+id: String +type: String +name: String +definition: String +metadata: Section +references: DataArray[] +positions: DataArray +extents: DataArray +units: String[] +features: Feature[] +sources: Source[]

NIX Model: Feature

- Used to attach additional data to a tag entity
- Example: attach a waveform to a spike

Feature
+id: String +link_type: LinkType +data: DataArray

NIX Model: Source

- Used to describe the provenance of other entities
- Can have child sources
- Example
 - Recording Channel
 - Electrode
 - Analysis

Source
+id: String +type: String +name: String +definition: String +metadata: Section +sources: Source[]

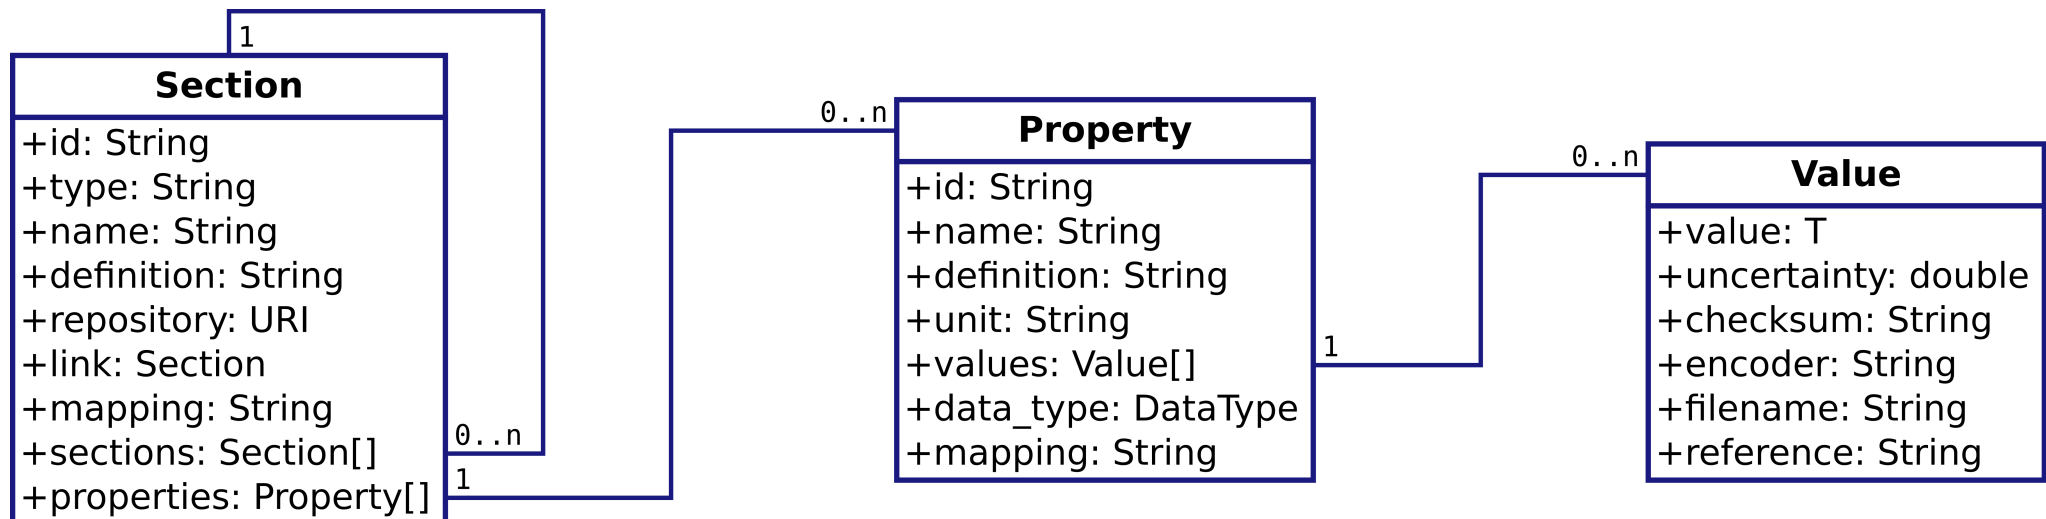
NIX Model: Block

- Group other entities such as sources, tags and data arrays

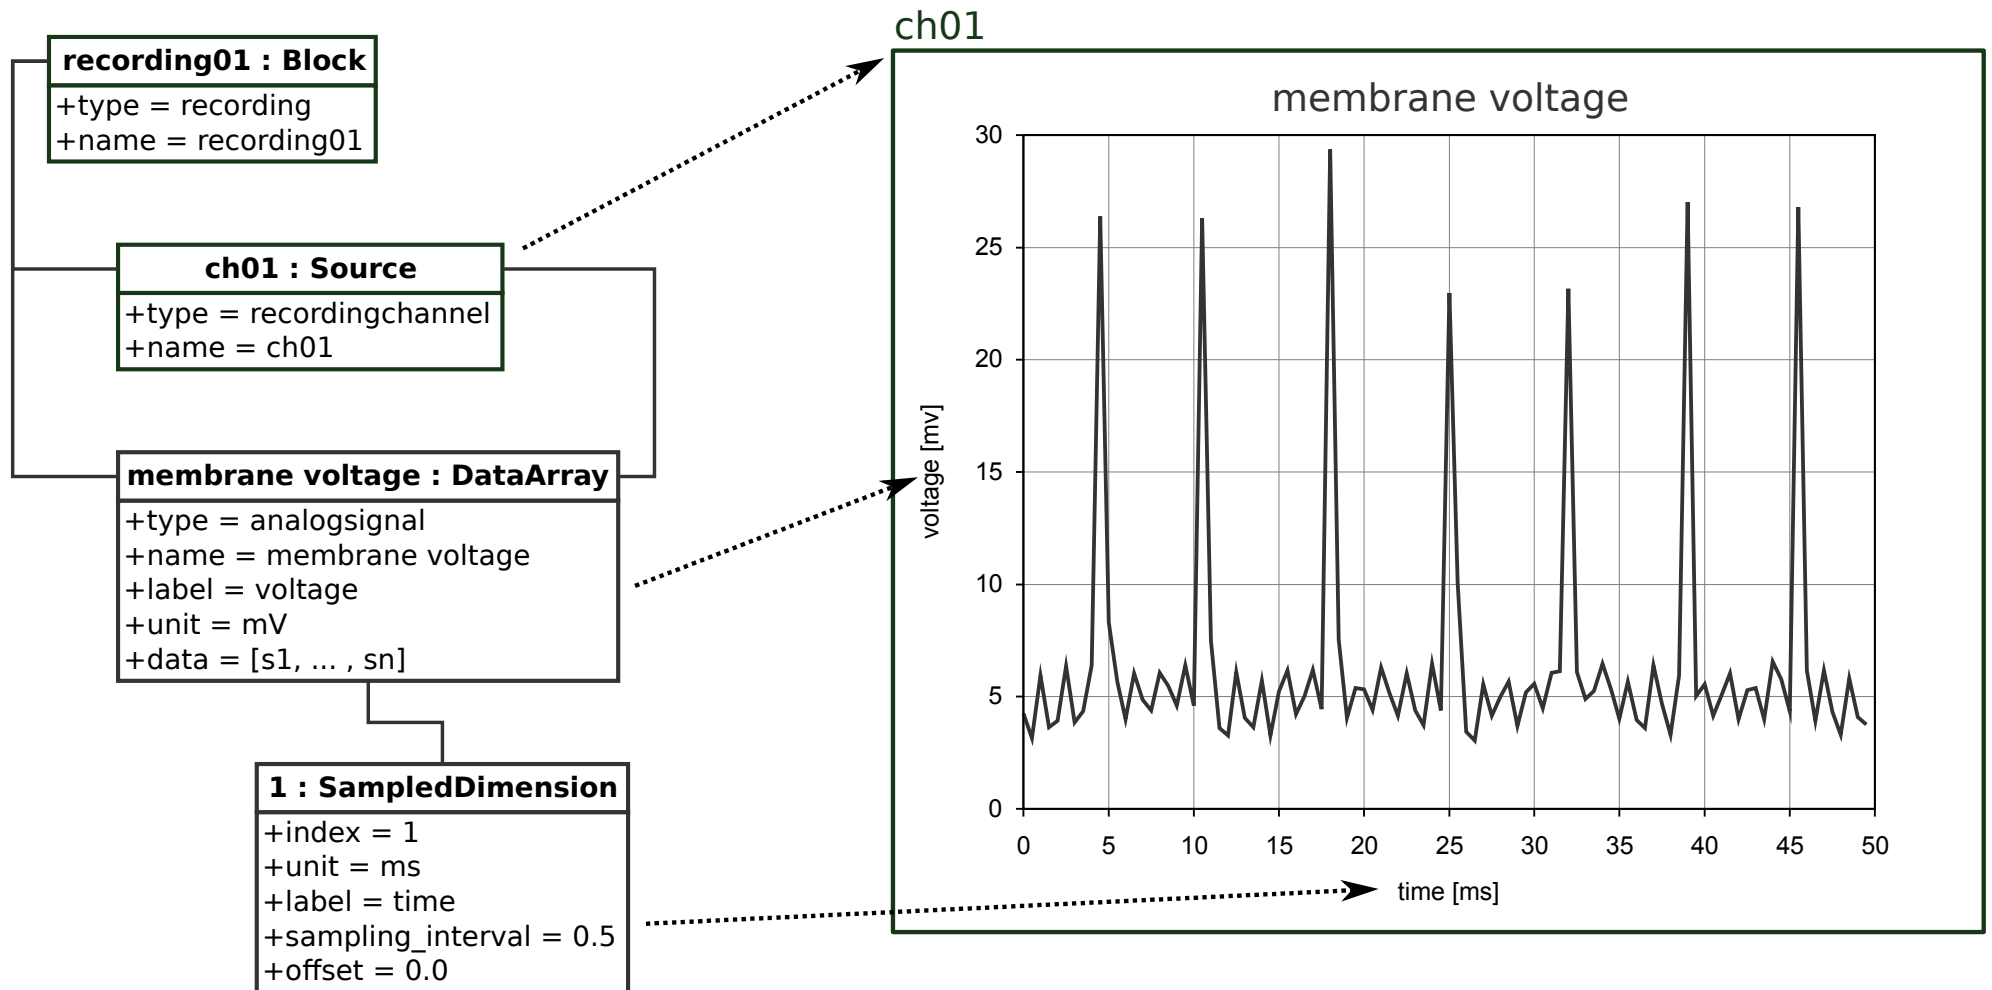
Block
+id: String +type: String +name: String +definition: String +metadata: Section +sources: Source[] +data_arrays: DataArray[] +multi_tags: MultiTag[] +tags: Tag[]

odML Model

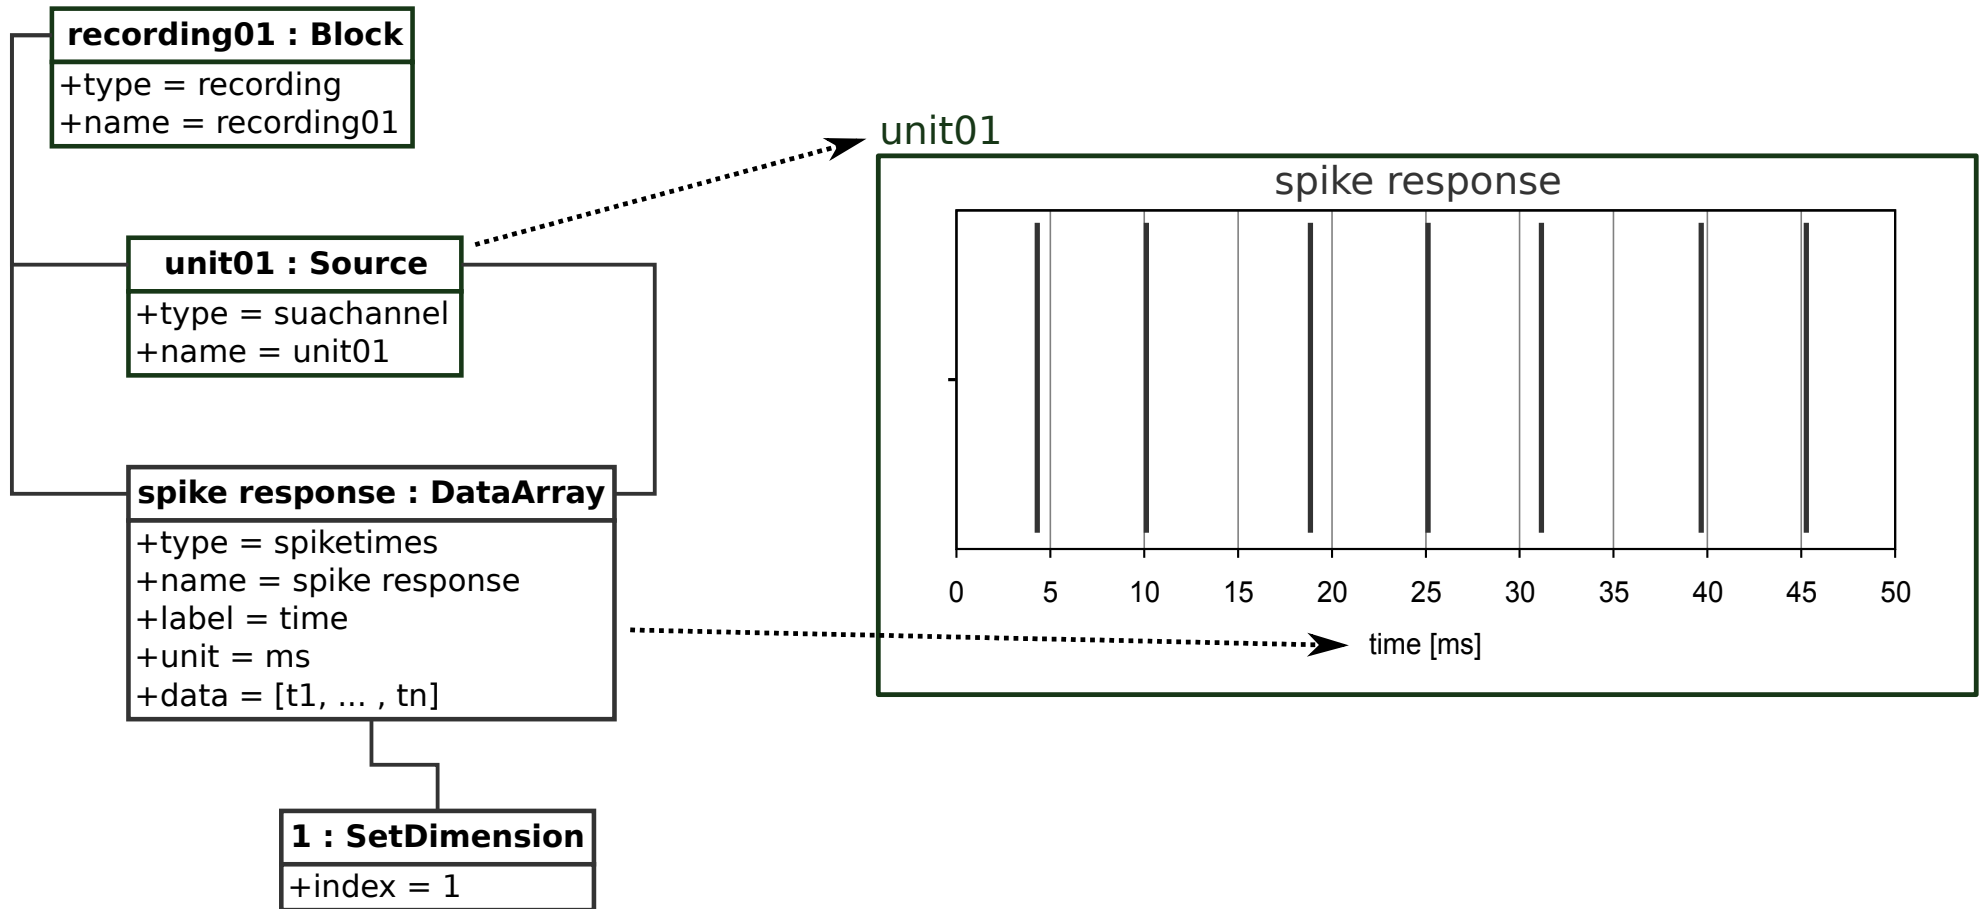
- Flexible model for metadata (Grewe et al. 2011)
- Stores hierarchically grouped key value pairs




Example: Analog Signal



Example: Spikes



Example: Signal and Spikes

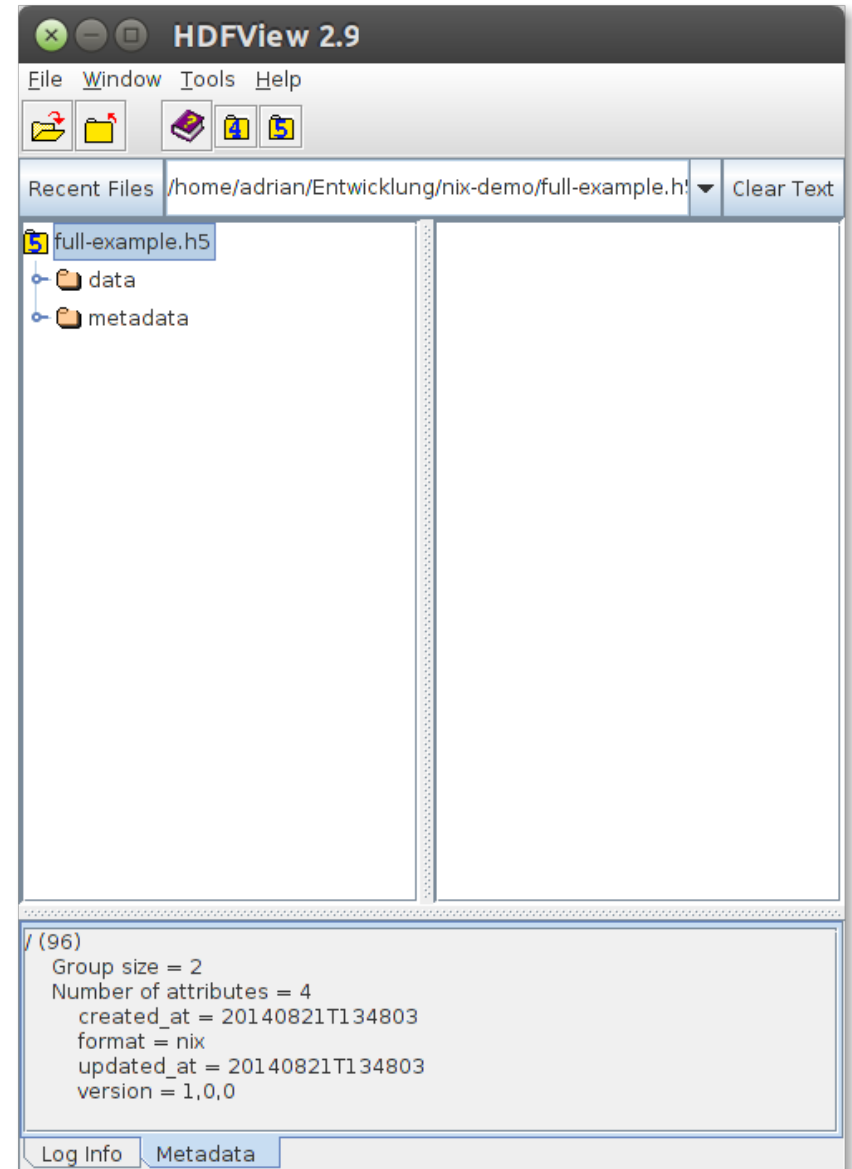


HDF5 Format: General Concepts

- Most entities are implemented as HDF5 groups
 - The entity name is the group name
- Each entity type resides at a specific location in the file
- Entity properties are HDF5 attributes
- Relationships are implemented as nested subgroups or as hard links to other groups

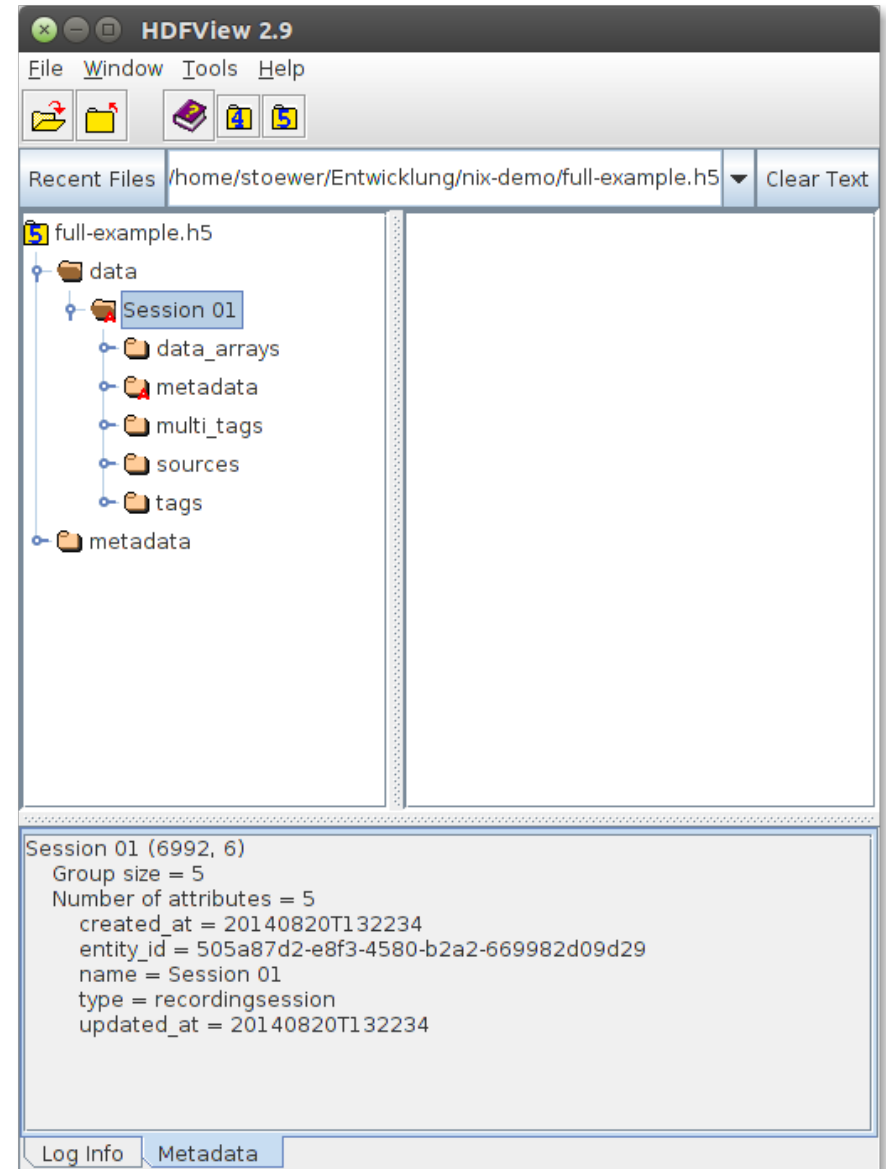
HDF5 Format: File Root

- File root attributes
 - Format hint
 - Format version
 - Update and creation time
- Subgroup *metadata*
 - Contains all odML root sections
- Subgroup *data*
 - Contains block entities



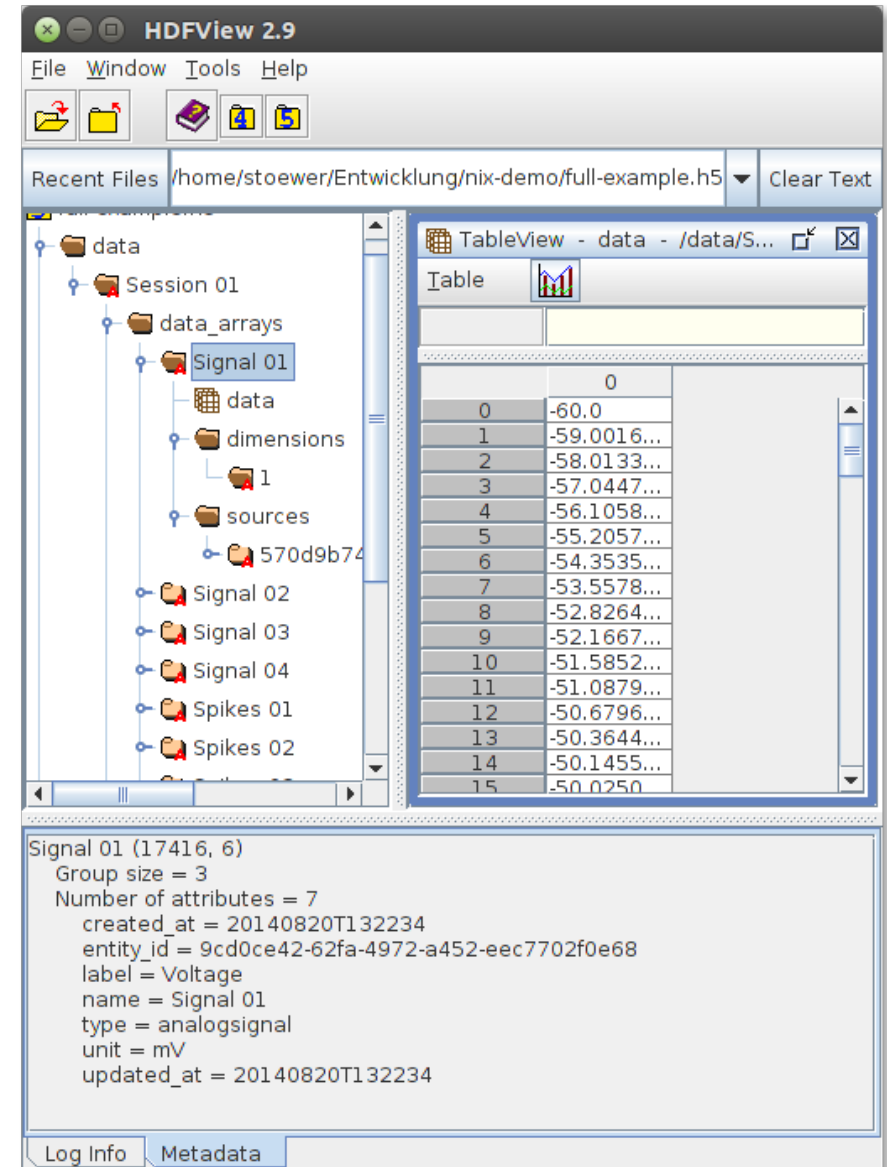
HDF5 Format: Block

- Subgroups for:
 - DataArrays
 - Tags
 - MultiTags
 - Sources



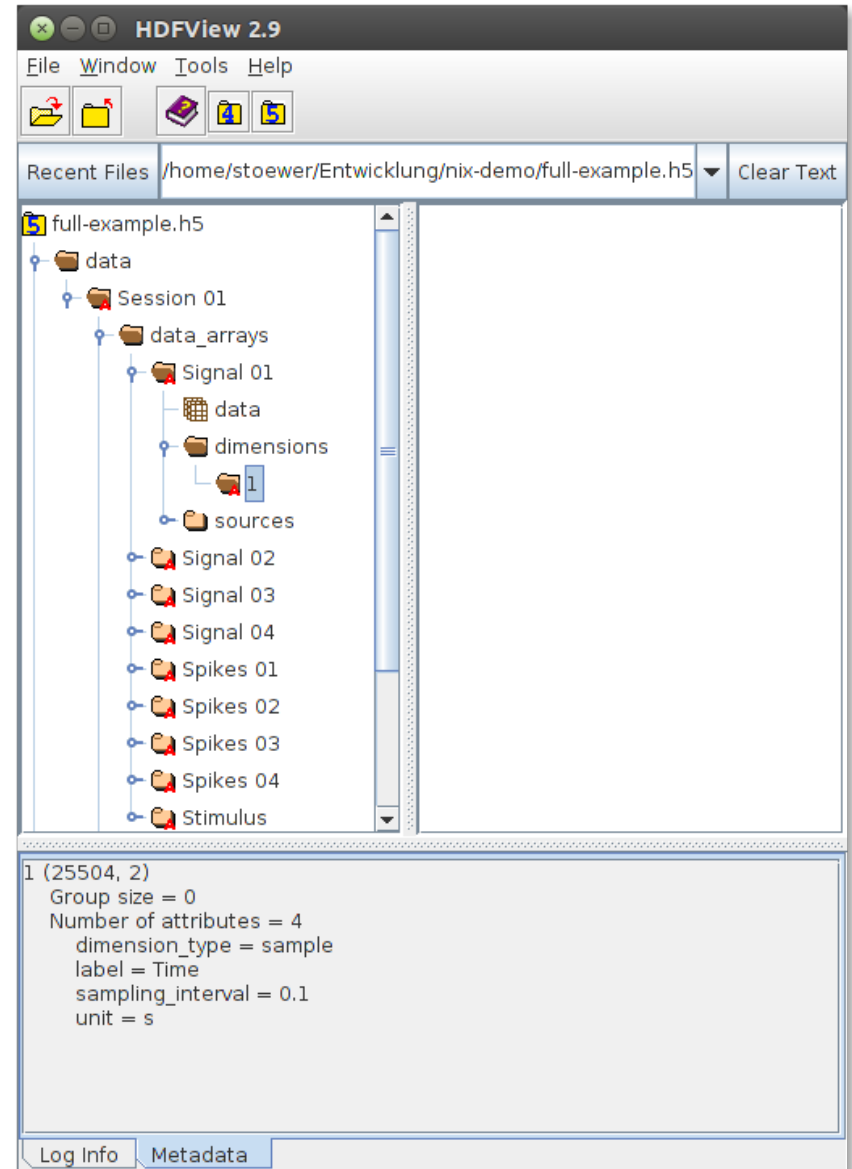
HDF5 Format: DataArray

- HDF5 dataset for the actual data
- Subgroup for dimension descriptors
- Subgroup for links to Source entities



HDF5 Format: Dimensions

- The group name is the dimension index
- RangeDimension entities contain an HDF5 dataset for the ticks



HDF5 Format: Tag

- HDF5 datasets for the position and extent vectors
- Subgroup for attached features
- Subgroup for links to referenced data

The screenshot shows the HDFView 2.9 application window. The main pane displays the file structure of 'full-example.h5'. The 'data' group is expanded, showing subgroups: 'Session 01', 'data_arrays', 'metadata', 'multi_tags', 'sources', 'tags', and 'Trial 01'. 'Trial 01' is selected, showing its contents: 'extent', 'features', 'position', 'references', and 'metadata'.

Two 'TableView' windows are open. The first, titled 'TableView - positi...', shows a table with two columns. The second, titled 'TableView - exten...', shows a table with two columns.

The bottom pane displays the metadata for 'Trial 01 (91368, 6)'. The metadata is as follows:

```
Trial 01 (91368, 6)
Group size = 4
Number of attributes = 5
  created_at = 20140820T132234
  entity_id = f00f288f-3d9b-404d-8b37-7322f52e7594
  name = Trial 01
  type = trial
  updated_at = 20140820T132234
```

The bottom status bar shows 'Log Info' and 'Metadata' tabs.

HDF5 Format: MultiTag

- Link to DataArray entities representing the positions and extents

The screenshot shows the HDFView 2.9 application window. The main pane displays a hierarchical tree of objects within a file named 'full-example.h5'. The tree structure is as follows:

- full-example.h5
 - data
 - Session 01
 - data_arrays
 - metadata
 - multi_tags
 - Tagged Spikes 0
 - positions
 - data (selected)
 - dimension
 - references
 - 9cd0ce42
 - Tagged Spikes 0
 - Tagged Spikes 0
 - Tagged Spikes 0

On the right, a 'TableView' window titled 'data - /data/S...' displays a table with 10 rows and 2 columns. The first column contains integers from 0 to 9, and the second column contains floating-point values from 0.0 to 9.0 in increments of 1.0.

0	0.0
1	1.0
2	2.0
3	3.0
4	4.0
5	5.0
6	6.0
7	7.0
8	8.0
9	9.0

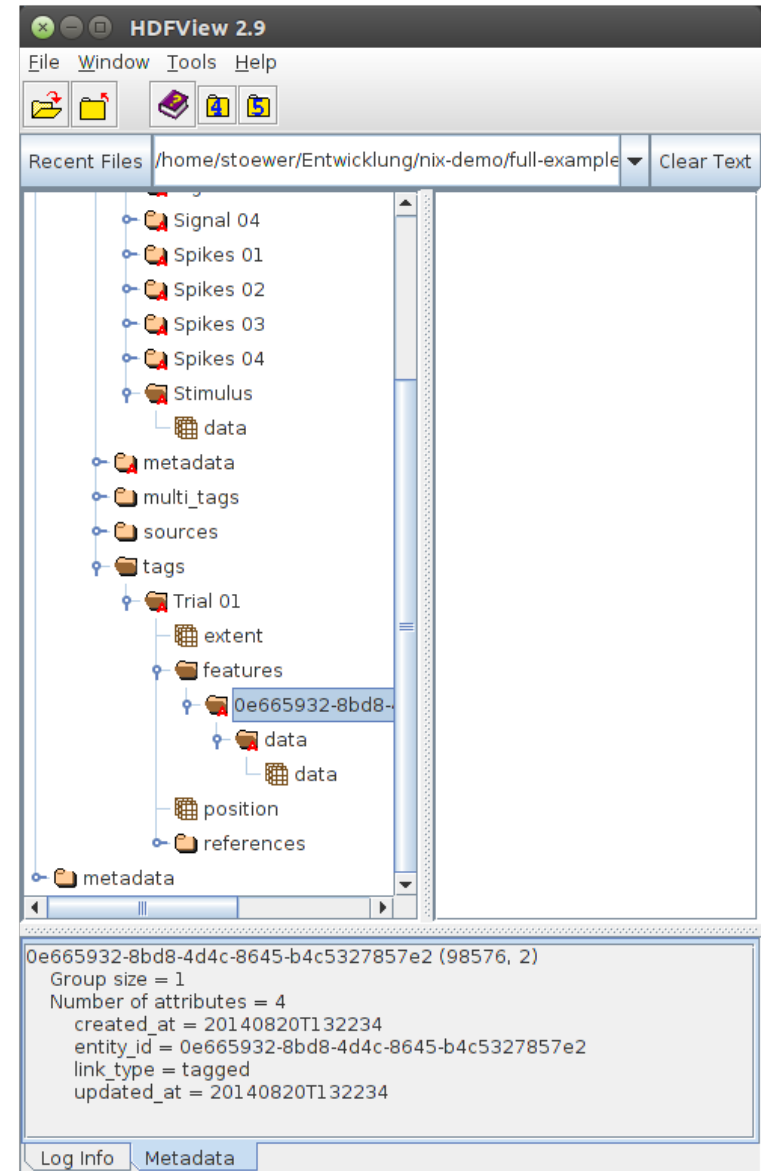
At the bottom of the main window, a metadata panel for the selected 'Tagged Spikes 01' object is displayed:

Tagged Spikes 01 (80352, 6)
Group size = 2
Number of attributes = 5
created_at = 20140820T132234
entity_id = 115bc502-aef0-454a-95ab-f1ae22731420
name = Tagged Spikes 01
type = spiketrain
updated_at = 20140820T132234

The bottom of the window has tabs for 'Log Info' and 'Metadata'.

HDF5 Format: Feature

- Link to the attached DataArray entity



HDF5 Format: Metadata

- Subgroup for child sections
- Subgroup for properties
- Properties and values are implemented as HDF5 datasets

The screenshot shows the HDFView 2.9 application window. The title bar reads 'HDFView 2.9'. The menu bar includes 'File', 'Window', 'Tools', and 'Help'. Below the menu bar is a toolbar with icons for file operations. The 'Recent Files' section shows the path '/home/stoewer/Entwicklung/nix-demo/full-example.h5'. The main panel displays a hierarchical tree structure of the HDF5 file's metadata. The tree is rooted at 'tags' and includes 'metadata', 'Session 01', 'properties', 'Experimenter', 'sections', 'Subject 01', 'Tetrode 01', and various datasets like 'Age', 'Name', 'Species', 'ElectrodeCount', and 'Manufacturer'. A 'TableView - Experimenter...' window is open, showing a table with columns 'value', 'uncertainty', and 're'. The table contains one row with the value 'Thomas ...' and an uncertainty of '0.0'. At the bottom of the main window, a text area displays metadata for 'Session 01 (100896, 6)', including group size, number of attributes, creation time, entity ID, name, type, and update time. The bottom status bar has tabs for 'Log Info' and 'Metadata'.

Session 01 (100896, 6)
Group size = 2
Number of attributes = 5
created_at = 20140820T132234
entity_id = 2803b955-47b0-44bb-95fd-a0041b92ebee
name = Session 01
type = recordingsession
updated_at = 20140820T132234

C++ API

- Reflects directly the concepts of the data model
- Multi compiler and multi platform
 - g++, clang, MSVC
 - Linux, OSX, Windows
- Designed to support multiple back-ends
- HDF5 back-end
 - Classes operate directly on the file
 - Allows partial reading and writing of data
- Memory back-end
 - Planned
- Hosted on GitHub
 - <https://github.com/G-Node/nix>

Python API

- Implemented as bindings to the native C++ API
- Numpy support
- Pythonic interface
 - Reading and writing works similar to h5py
- Also hosted on GitHub
 - <https://github.com/G-Node/nixpy>

License

- NIX is open source
- Published under BSD License
 - Allows use and modification of the source code
 - Can be used in commercial non open source products
 - Even modified versions of NIX can be distributed without source code

Project Status

- File format is stable
- API is stable
- Beta 1 release in preparation
 - Packages for Debian via Launchpad
<https://launchpad.net/~gnode/+archive/ubuntu/nix>
 - Python package on PyPi
 - Compiled version for Windows

Tools using NIX

- RELACS
 - A user interface for data acquisition
 - It is controllable and customizable to specific experimental setup by C++ based plugins
- Guppy
 - A small video recorder tool
 - Stores videos grabbed from open CV devices either as avi or to the NIX file format

Outlook

- Improve model, format and API
- Implement other back-ends
 - In-memory back-end
- More language bindings
 - Java
 - Matlab
- Think about RDF support
 - Metadata could be stored directly as triples
 - NIX type could be interpreted as `rdf:type`
- Make NIX a community project

Acknowledgments

Programming and Data Model:

Adrian Stoewer
Christian Kellner
Jan Grewe
Balint Morvai
Andrey Sobolev

Project Lead:

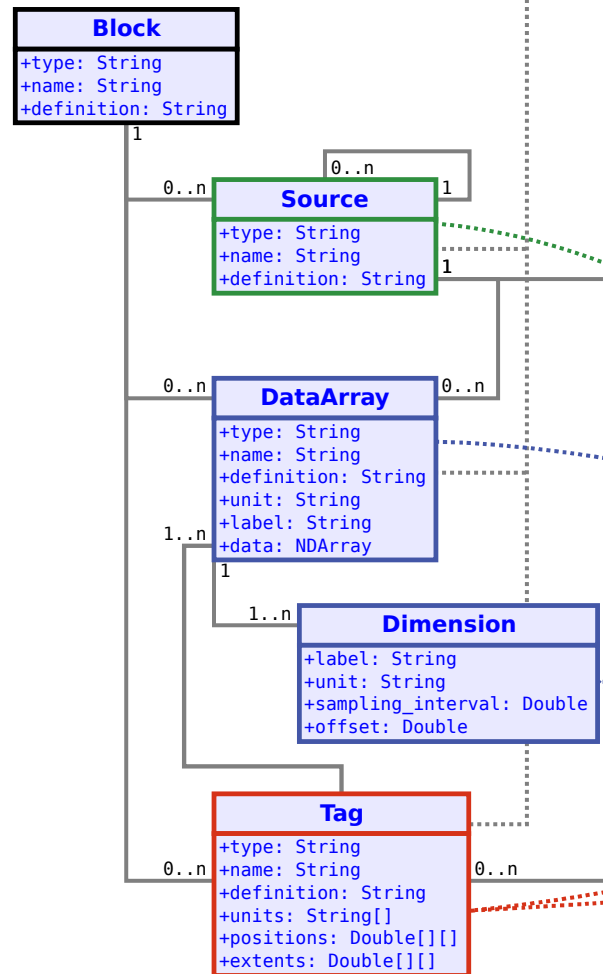
Thomas Wachtler
Jan Grewe

Funding:

BMBF, grant 01GQ1302

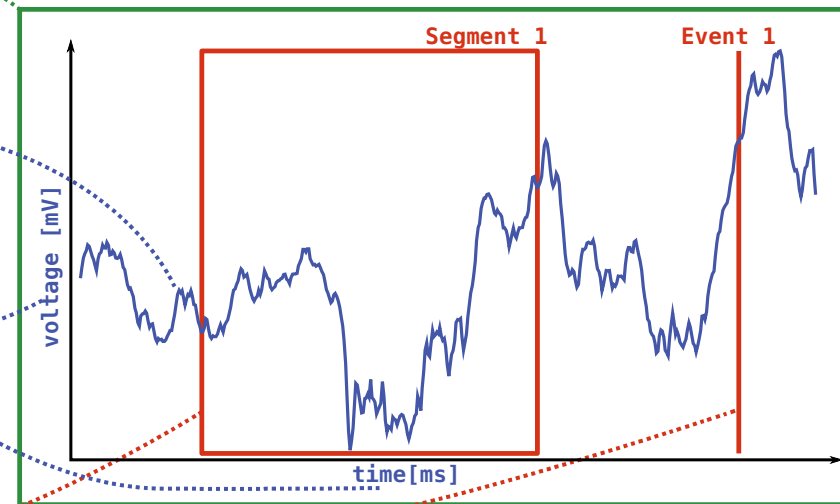
Metadata Model (odML)

Data Model



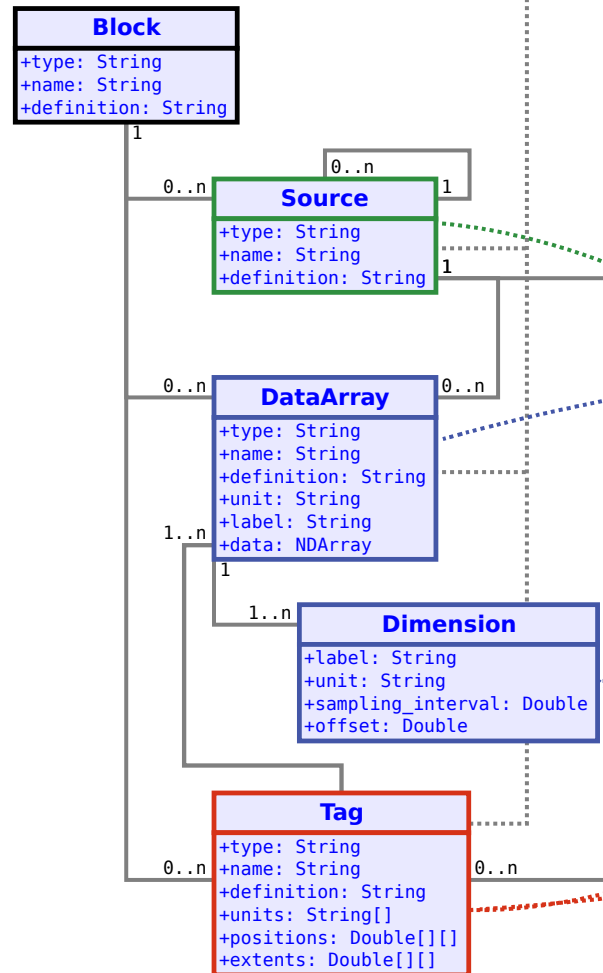
Example Data

Recordingchannel 2



Metadata Model (odML)

Data Model



Example Data

Laser Scanning Microscope

